


Training Neural Networks

Mini-batch SGD

- Loop:
1. Sample a batch of data
 2. Forward prop it through the graph, get loss
 3. Backprop to calculate the gradients
 4. Update the parameters using the gradients


Activation Functions

Sigmoid $\sigma(x) = \frac{1}{(1+e^{-x})}$




- saturates "firing rate" of a neuron
- Saturated neurons kill the gradients
- Sigmoid outputs are not zero-centered
- $\exp()$ is a bit compute expensive

tanh(x)



- squashes numbers to $[-1, 1]$
- zero centered
- still kills gradients when saturated

ReLU



- computes $f(x) = \max(0, x)$
- does not saturate
- very computationally efficient
- converges much faster than Sigmoid, tanh
- actually more biologically plausible than Sigmoid.
- Not zero-centered

Maxout "Neuron" $\max(w_1^T x + b_1, w_2^T x + b_2)$

- not the form of dot product : nonlinearity
- generalizes ReLU and leaky ReLU
- Linear regime, does not saturate/die
- doubles the number of parameters/neuron

결론: 실전에서 ReLU가 가장 좋은 성능

Data Preprocessing

in practice for images : center only

not common to normalize variance,
to do PCA or whitening

Weight Initialization

1st idea : Small random numbers


(가우시안 분포 0 평균 1e-2 표준편차)

Deep layer \rightarrow all activations become zero

"Xavier Initialization" "He"

Batch Normalization (for unit gaussian activations) Variance를 줄이기 위해

1. compute the empirical mean and variance independently for each dimension



2. Normalize

$$\hat{x}^{(k)} = \frac{x^{(k)} - E[x^{(k)}]}{\sqrt{\text{Var}[x^{(k)}]}}$$

- ∴ Improves gradient flow through the network
- Allows higher learning rates
- Reduces the strong dependence on initialization
- Regularization Effect

Z-score Normalization : $(X - \mu) / \sigma$

Random Search > Grid Search

Loss Curve visualization ∴ learning rate 조정